

Week 1: Overview and basic configurations

Step 1: Choose a suitable cloud provider and set up a Spark shell environment

Step 2: Configure the necessary dependencies

Step 3: Execute basic Spark commands to make sure Spark is ready

Ecommerce Insights

Setting Up Environment

M ↩ ↶

```
1 from pyspark.sql import SparkSession
2 from pyspark.sql.functions import *
3 from pyspark.sql.types import *
4
5 import datetime as dt
```

✓ - Command executed in 149 ms on 8:33:02 AM, 9/05/24

Step 4: Use README.md for details, instructions, and commands

Week 2: Data ingestion

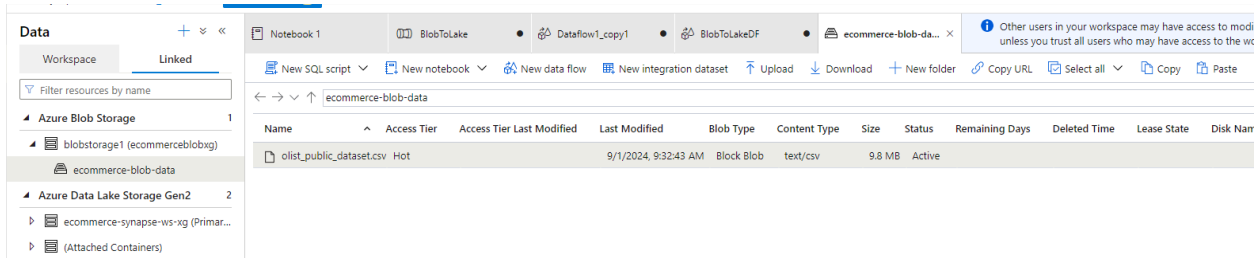
Step 1: Upload the entire data into Hive from CSV using cloud provider cluster setup

- 1. Log in to PuTTY with the username “hadoop”
- 2. Enter the command given below:

Command: hive

- 3. Create a database
- 4. Create a table with all the relevant details

Step 2: Create a bucket (Azure Blob) and upload the csv file



Step 4: Create a new directory in HDFS(Data Lake) and copy the data from Hive into HDFS

Data | Workspace | **Linked**

Filter resources by name

- Azure Blob Storage 1
 - blobstorage1 (ecommerceblobxg)
 - ecommerce-blob-data
- Azure Data Lake Storage Gen2 2
 - ecommerce-synapse-ws-xg (Primar...)
 - (Attached Containers)
- Integration datasets 2
 - ecommerceCSVBlobDS
 - ecommerceCSVLakeDS

Notebook 1 | BlobToLake | Dataflow1_copy1 | BlobToLakeDF | ecommerce-blob-data | ecommerceCSVBlobDS

DelimitedText
ecommerceCSVBlobDS

Connection | Schema | Parameters

Linked service * | blobstorage1 | Test connection | Edit | New | Learn more

Integration runtime * | AutoResolveIntegrationRuntime | Edit

File path * | ecommerce-blob-data | / Directory | / olist_public_dataset.csv | Browse | Preview data | Det

Compression type | Select...

Column delimiter | Comma (,)

Row delimiter | Default \n\r\n

Data | Workspace | **Linked**

Filter resources by name

- Azure Blob Storage 1
 - blobstorage1 (ecommerceblobxg)
 - ecommerce-blob-data
- Azure Data Lake Storage Gen2 2
 - ecommerce-synapse-ws-xg (Primar...)
 - (Attached Containers)
- Integration datasets 2
 - ecommerceCSVBlobDS
 - ecommerceCSVLakeDS

Notebook 1 | BlobToLake | Dataflow1_copy1 | BlobToLakeDF | ecommerce-blob-data | ecommerceCSV...

DelimitedText
ecommerceCSVLakeDS

Connection | Schema | Parameters

Linked service * | ecommerce-synapse-ws-xg-Worksp... | Test connection | Edit | New | Learn more

Integration runtime * | AutoResolveIntegrationRuntime | Edit

File path | ecommercecontainerxg | / Directory | / File name | Browse | Previe

Compression type | Select...

Column delimiter | Comma (,)

Synapse live | Validate all | Publish all

Develop | Filter resources by name

- SQL scripts 1
- Notebooks 1
 - Notebook 1
- Data flows 2
 - BlobToLakeDF
 - Dataflow1_copy1

Notebook 1 | BlobToLake | Dataflow1_copy1 | BlobToLakeDF | ecommerce-blob-data | ecommerceCSVBlob... | ecommerceCSVLake...

Validate | Data flow debug

Columns: 15 total

Source settings | Source options | **Projection** | Optimize | Inspect | Data preview

Define default format | Detect data type | Import projection | Reset schema

Column name	Type	Format
Id	integer	Specify format
order_status	string	Specify format
order_products_value	float	Specify format
order_freight_value	float	Specify format
order_items_qty	short	Specify format
customer_city	string	Specify format
customer_state	string	Specify format
customer_zip_code_prefix	integer	Specify format
product_name_length	integer	Specify format
product_description_length	integer	Specify format
product_photos_qty	integer	Specify format
review_score	integer	Specify format
order_purchase_timestamp	timestamp	dd/MM/yy HH:mm
order_approved_at	timestamp	dd/MM/yy HH:mm
order_delivered_customer_date	timestamp	dd/MM/yy HH:mm

Home

SQL pools

Apache Spark pools

Data Explorer pools (preview)

Activities

SQL requests

KQL requests

Apache Spark applications

Data flow debug

Integration

Pipeline runs

Trigger runs

Integration runtimes

Link connections

All pipeline runs > Pipeline 1 - Activity runs

Refresh

Update pipeline

List

Gantt

Data flow

Data flow1

Activity runs

Pipeline run ID e85650e2-20bf-44c2-bcd9-38bf8211462f

All status

Showing 1 - 1 of 1 items

Activity name	Activity status	Activity type	Run start	Duration	Integration runtime	User properties
Data flow1	Succeeded	Data flow	9/1/2024, 10:28:55 AM	1m 9s	AutoResolveIntegrationRu	

Step 3: Load the data from the bucket into the Hive table(Lake Database)

Create external table from data lake

External table details

Select the storage location where the files containing the data is staged. Currently Azure Data Lake Storage (ADLS) Gen2 and Azure Blob Storage are supported. [Learn more](#)

External table name *

Linked service * ⓘ

Input file or folder * ⓘ

New external table

Source file format settings

Specify the format and layout of your data. [Learn more](#)

File path

ecommercecontainerxg/part-00000-660c7ed8-6c9c-4961-a361-7e1a973e413b-c000.csv

[Preview Data](#)

File type

CSV

Field terminator ⓘ

Default (comma ,)

☐ Edit

First row

☒ Infer column names ⓘ

String delimiter ⓘ

Default (Empty string)

☐ Edit

Use default type ⓘ

Default type (true,false)

Max string length * ⓘ

4000

The screenshot shows the Azure Data Explorer interface. On the left, the 'Data' pane displays a tree view of resources under 'Lake database' > 'Database1' > 'Tables' > 'ecommerce_data'. The columns listed are: id (long), order_status (string), order_products_value (double), order_freight_value (double), order_items_qty (long), customer_city (string), customer_state (string), customer_zip_code_prefix (string), product_name_length (long), product_description_length (long), product_photos_qty (long), review_score (long), order_purchase_timestamp (timestamp), order_approved_at (timestamp), and order_delivered_customer_time (timestamp).

The main pane shows the 'ecommerce_data' table details. The 'General' tab is active, displaying the table name 'ecommerce_data' and a description field labeled 'Enter a description'. Below this, the 'Storage settings for table' section is visible.

Step 5: Check if the data has been successfully loaded in the HDFS path

The screenshot shows the Azure Data Explorer interface. On the left, the 'Data' pane displays a tree view of resources under 'Azure Blob Storage' > 'ecommerce-blob-storage' > 'ecommercecontainer'. The main pane shows a table of files and folders within the 'ecommercecontainer' storage path.

Name	Last Modified	Content Type
synapse	9/1/2024, 10:16:23 AM	Folder
_SUCCESS	9/1/2024, 10:29:14 AM	
part-00000-660c7ed8-6c9c-4961-a361-7e1a973e413b-c000.csv	9/1/2024, 10:29:12 AM	

Week 3: Data streaming

Step 1: Connect to Spark shell with all the dependencies (Hive, Hadoop, and HDFS).

1. Create Schema of the CSV files

Creating Schema

✓

```
1  ecommerce_schema = StructType([
2      StructField("id", IntegerType(), False),
3      StructField("order_status", StringType(), True),
4      StructField("order_products_value", FloatType(), True),
5      StructField("order_freight_value", FloatType(), True),
6      StructField("order_items_qty", IntegerType(), True),
7      StructField("customer_city", StringType(), True),
8      StructField("customer_state", StringType(), True),
9      StructField("customer_zip_code_prefix", IntegerType(), True),
10     StructField("product_name_lenght", IntegerType(), True),
11     StructField("product_description_lenght", IntegerType(), True),
12     StructField("product_photos_qty", IntegerType(), True),
13     StructField("review_score", IntegerType(), True),
14     StructField("order_purchase_timestamp", TimestampType(), True),
15     StructField("order_approved_at", TimestampType(), True),
16     StructField("order_delivered_customer_date", TimestampType(), True)
17 ])
```

[7] ✓ - Command executed in 152 ms on 8:33:41 AM, 9/05/24

...

2. Create a Spark session(spark session is already up but this is what it would look like)

▷ ✓

```
1  ecommerce_session = SparkSession.builder.appName("Ecommerce").getOrCreate()
```

- Add Object Storage Service details as per the Cloud provider
- Add all variables to your environment as they contain sensitive data

Step 2: Read the CSV file and convert the file to a data frame

Creating Dataframe

MI

```
1 ecommerce_df = spark.read.format("csv").option("header", "True").schema(ecommerce_schema).load("abfss://ecommerce@storage.azure.com/ecommerce.csv")
2 print((ecommerce_df.count(), len(ecommerce_df.columns)))
3 ecommerce_df.printSchema()
```

[8] ✓ - Command executed in 14 sec 648 ms on 8:33:58 AM, 9/05/24

```
*** (100000, 15)
root
 |-- id: integer (nullable = true)
 |-- order_status: string (nullable = true)
 |-- order_products_value: float (nullable = true)
 |-- order_freight_value: float (nullable = true)
 |-- order_items_qty: integer (nullable = true)
 |-- customer_city: string (nullable = true)
 |-- customer_state: string (nullable = true)
 |-- customer_zip_code_prefix: integer (nullable = true)
 |-- product_name_lenght: integer (nullable = true)
 |-- product_description_lenght: integer (nullable = true)
 |-- product_photos_qty: integer (nullable = true)
 |-- review_score: integer (nullable = true)
 |-- order_purchase_timestamp: timestamp (nullable = true)
 |-- order_approved_at: timestamp (nullable = true)
 |-- order_delivered_customer_date: timestamp (nullable = true)
```

Step 3: Convert “order_purchase_timestamp” to week and day using UDF

Setting up Dataframe

Creating Date Columns

```
1 ecommerce_df = ecommerce_df.withColumn("order_purchase_as_date",to_date(col("order_purchase_timestamp"),'dd/MM/yy'))
2 ecommerce_df = ecommerce_df.withColumn("order_approved_at_as_date",to_date(col("order_approved_at"),'dd/MM/yy'))
3 ecommerce_df = ecommerce_df.withColumn("order_delivery_customer_date_as_date",to_date(col("order_delivered_customer_date"),'dd/MM/yy'))
```

✓ - Command executed in 146 ms on 8:34:00 AM, 9/05/24

```
1 ecommerce_df.printSchema()
2 display(ecommerce_df)
```

[10] ✓ - Command executed in 3 sec 990 ms on 8:34:05 AM, 9/05/24

```
... root
|-- id: integer (nullable = true)
|-- order_status: string (nullable = true)
|-- order_products_value: float (nullable = true)
|-- order_freight_value: float (nullable = true)
|-- order_items_qty: integer (nullable = true)
|-- customer_city: string (nullable = true)
|-- customer_state: string (nullable = true)
|-- customer_zip_code_prefix: integer (nullable = true)
|-- product_name_lenght: integer (nullable = true)
|-- product_description_lenght: integer (nullable = true)
|-- product_photos_qty: integer (nullable = true)
|-- review_score: integer (nullable = true)
|-- order_purchase_timestamp: timestamp (nullable = true)
|-- order_approved_at: timestamp (nullable = true)
|-- order_delivered_customer_date: timestamp (nullable = true)
|-- order_purchase_as_date: date (nullable = true)
|-- order_approved_at_as_date: date (nullable = true)
|-- order_delivery_customer_date_as_date: date (nullable = true)
```

Creating a Day of Week Column

```
1 def dayAsString(day_int):
2     day_of_week_string = {
3         1: "Monday",
4         2: "Tuesday",
5         3: "Wednesday",
6         4: "Thursday",
7         5: "Friday",
8         6: "Saturday",
9         7: "Sunday"
10    }
11    return day_of_week_string[day_int]
```

[11] ✓ - Command executed in 157 ms on 8:34:05 AM, 9/05/24

```
1 day_as_string_udf = udf(lambda x: dayAsString(x), StringType())
```

[12] ✓ - Command executed in 198 ms on 8:34:05 AM, 9/05/24

```
1 ecommerce_df = ecommerce_df.withColumn("day_of_week",day_as_string_udf(dayofweek(col("order_purchase_as_date"))))
2 display(ecommerce_df)
```

[13] ✓ - Command executed in 10 sec 676 ms on 8:34:16 AM, 9/05/24

```
1 ecommerce_df.select("day_of_week").distinct().show()
```

[14] ✓ - Command executed in 1 sec 832 ms on 8:34:18 AM, 9/05/24

```
...
+-----+
|day_of_week|
+-----+
| Wednesday|
| Tuesday  |
| Friday   |
| Thursday |
| Saturday |
| Monday   |
| Sunday   |
+-----+
```

Creating a Week of Month Column

```
1 def weekOfMonth(date1):
2     month = date1.month
3     year = date1.year
4     month_start = dt.date(year, month, 1)
5     if (dt.date(year, 1, 1) == month_start):
6         week_of_year_month_start = 1
7     else:
8         week_of_year_month_start = month_start.isocalendar()[1]
9     week_of_year_date1 = date1.isocalendar()[1]
10    return week_of_year_date1 - week_of_year_month_start + 1
```

[15] ✓ - Command executed in 152 ms on 8:34:18 AM, 9/05/24

```
1 weekOfMonth_udf = udf(lambda x: weekOfMonth(x), IntegerType())
```

[16] ✓ - Command executed in 155 ms on 8:34:19 AM, 9/05/24

```
1 ecommerce_df = ecommerce_df.withColumn("week_of_month", weekOfMonth_udf(col("order_purchase_as_date")))
2 display(ecommerce_df)
```

[17] ✓ - Command executed in 1 sec 925 ms on 8:34:21 AM, 9/05/24


```
1 ecommerce_df.select("week_of_month").distinct().show()
```

✓ - Command executed in 1 sec 812 ms on 8:34:23 AM, 9/05/24

```
+-----+
|week_of_month|
+-----+
|           1|
|           6|
|           3|
|           5|
|           4|
|           2|
+-----+
```

```
1 ecommerce_df.printSchema()
```

[19] ✓ - Command executed in 143 ms on 8:34:23 AM, 9/05/24

```
... root
  |-- id: integer (nullable = true)
  |-- order_status: string (nullable = true)
  |-- order_products_value: float (nullable = true)
  |-- order_freight_value: float (nullable = true)
  |-- order_items_qty: integer (nullable = true)
  |-- customer_city: string (nullable = true)
  |-- customer_state: string (nullable = true)
  |-- customer_zip_code_prefix: integer (nullable = true)
  |-- product_name_lenght: integer (nullable = true)
  |-- product_description_lenght: integer (nullable = true)
  |-- product_photos_qty: integer (nullable = true)
  |-- review_score: integer (nullable = true)
  |-- order_purchase_timestamp: timestamp (nullable = true)
  |-- order_approved_at: timestamp (nullable = true)
  |-- order_delivered_customer_date: timestamp (nullable = true)
  |-- order_purchase_as_date: date (nullable = true)
  |-- order_approved_at_as_date: date (nullable = true)
  |-- order_delivery_customer_date_as_date: date (nullable = true)
  |-- day_of_week: string (nullable = true)
  |-- week_of_month: integer (nullable = true)
```

Step 4: Calculate the following data:

1. Total sales and order distribution per day and week for each city
2. Total sales and order distribution per day and week for each state
3. Average review score, average freight value, average order approval, and delivery time
4. The freight charges per city and total freight charges

Getting Ecommerce Results

Making Dictionary to Store Results

```
1 ecommerce_datasets = {}
```

✓ - Command executed in 176 ms on 8:34:23 AM, 9/05/24

Insight on Sales

```
1 total_sales_df = ecommerce_df.agg(round(sum("order_products_value"),2).alias("total_sales"))
2 total_sales_df.show()
3
4 ecommerce_datasets["total_sales"] = total_sales_df
```

✓ - Command executed in 1 sec 58 ms on 8:34:24 AM, 9/05/24

```
+-----+
| total_sales|
+-----+
|1.284147698E7|
+-----+
```

Sales by Day

```
1 total_sales_by_day_df = ecommerce_df.groupBy("day_of_week").agg(round(sum("order_products_value"),2).alias("total_sales")).orderBy("total_sales")
2 total_sales_by_day_df.show()
3
4 ecommerce_datasets["total_sales_by_day"] = total_sales_by_day_df
```

[51] ✓ - Command executed in 1 sec 890 ms on 8:50:44 AM, 9/05/24

```
+-----+-----+
|day_of_week|total_sales|
+-----+-----+
| Sunday| 1406302.11|
| Monday| 1530648.75|
| Saturday| 1811400.14|
| Friday| 1922774.44|
| Thursday| 1989730.72|
| Wednesday| 2077208.79|
| Tuesday| 2103412.03|
+-----+-----+
```

```
1 total_sales_by_day_and_city_df = ecommerce_df.groupBy("day_of_week", "customer_city").agg(round(sum("order_products_value"),2)\
2 .alias("total_sales")).orderBy("customer_city")
3 total_sales_by_day_and_city_df.show()
4
5 ecommerce_datasets["total_sales_by_day_and_city"] = total_sales_by_day_and_city_df
```

[52] ✓ - Command executed in 1 sec 845 ms on 8:50:46 AM, 9/05/24

```

.. +-----+-----+-----+
|day_of_week|customer_city|total_sales|
+-----+-----+-----+
| Friday|ALMIRANTE TAMANDA...| 49.9|
| Tuesday|ALMIRANTE TAMANDA...| 99.99|
| Sunday|ALTA FLORESTA D'O...| 708.99|
| Tuesday|ALTO ALEGRE DOS P...| 299.0|
| Sunday|ALTO ALEGRE DOS P...| 299.0|
| Monday|ALTO ALEGRE DOS P...| 314.99|
| Friday| ALVORADA D'OESTE| 328.0|
| Tuesday| ALVORADA D'OESTE| 359.98|
| Saturday| Abadia dos Dourados| 39.9|
| Tuesday| Abadia dos Dourados| 319.0|
| Tuesday| Abadiania| 68.9|
| Friday| Abadiania| 949.99|
| Thursday| Abaete| 398.79|
| Tuesday| Abaete| 449.0|
| Wednesday| Abaete| 321.6|
| Friday| Abaete| 56.99|
| Wednesday| Abaetetuba| 435.41|
| Monday| Abaetetuba| 115.99|
| Saturday| Abaetetuba| 1574.8|
| Friday| Abaetetuba| 164.89|
+-----+-----+-----+
only showing top 20 rows

```

```

✓ 1 total_sales_by_day_and_state_df = ecommerce_df.groupBy("day_of_week","customer_state").agg(round(sum("order_products_value"),2)\
2 |.alias("total_sales")).orderBy("customer_state")
3 total_sales_by_day_and_state_df.show()
4
5 ecommerce datasets["total sales by day and state"] = total sales by day and state df

```

```

+-----+-----+-----+
|day_of_week|customer_state|total_sales|
+-----+-----+-----+
| Thursday| AC| 2862.77|
| Monday| AC| 627.57|
| Sunday| AC| 853.49|
| Friday| AC| 2542.04|
| Saturday| AC| 3987.15|
| Wednesday| AC| 3478.97|
| Tuesday| AC| 2548.83|
| Saturday| AL| 7661.65|
| Sunday| AL| 8837.77|
| Monday| AL| 13089.56|
| Tuesday| AL| 14769.35|
| Wednesday| AL| 10002.48|
| Friday| AL| 11160.45|
| Thursday| AL| 12779.28|
| Sunday| AM| 2519.56|
| Friday| AM| 2063.87|
| Wednesday| AM| 2145.55|
| Saturday| AM| 3192.41|
| Monday| AM| 3949.5|
| Thursday| AM| 2505.38|
+-----+-----+-----+
only showing top 20 rows

```

Sales by Week

```
1 total_sales_by_week_df = ecommerce_df.groupBy("week_of_month").agg(round(sum("order_products_value"),2).alias("total_sales")).orderBy("week_of_month")
2 total_sales_by_week_df.show()
3
4 ecommerce_datasets["total_sales_by_week"] = total_sales_by_week_df
```

✓ - Command executed in 1 sec 77 ms on 8:50:49 AM, 9/05/24

week_of_month	total_sales
1	1542155.29
2	3089498.12
3	2957336.5
4	3027674.28
5	2058842.95
6	165969.84

```
1 total_sales_by_week_and_city_df = ecommerce_df.groupBy("week_of_month","customer_city").agg(round(sum("order_products_value"),2)\
2 |.alias("total_sales")).orderBy("customer_city")
3 total_sales_by_week_and_city_df.show()
4
5 ecommerce_datasets["total_sales_by_week_and_city"] = total_sales_by_week_and_city_df
```

✓ - Command executed in 1 sec 782 ms on 8:50:51 AM, 9/05/24

week_of_month	customer_city	total_sales
2	ALMIRANTE TAMANDA...	99.99
5	ALMIRANTE TAMANDA...	49.9
4	ALTA FLORESTA D'O...	349.99
1	ALTA FLORESTA D'O...	359.0
3	ALTO ALEGRE DOS P...	299.0
4	ALTO ALEGRE DOS P...	613.99
1	ALVORADA D'OESTE	328.0
5	ALVORADA D'OESTE	359.98
3	Abadia dos Dourados	358.9
5	Abadiania	68.9
1	Abadiania	949.99
5	Abaete	515.7
4	Abaete	176.89
2	Abaete	254.9
1	Abaete	69.99
3	Abaete	208.9
3	Abaetetuba	2410.96
1	Abaetetuba	134.99
2	Abaetetuba	305.51
4	Abaetetuba	63.8

nly showing top 20 rows

```
1 total_sales_by_week_and_state_df = ecommerce_df.groupBy("week_of_month","customer_state").agg(round(sum("order_products_value"),2)\
2 |.alias("total_sales")).orderBy("customer_state")
3 total_sales_by_week_and_state_df.show()
4
5 ecommerce_datasets["total_sales_by_week_and_state"] = total_sales_by_week_and_state_df
```

week_of_month	customer_state	total_sales
1	AC	3787.2
4	AC	2092.61
5	AC	4023.55
3	AC	4034.82
6	AC	699.0
2	AC	2263.64
6	AL	1057.29
4	AL	18700.29
1	AL	7343.29
3	AL	17451.02
2	AL	19961.02
5	AL	13787.63
4	AM	5775.42
3	AM	5308.23
5	AM	1554.8
2	AM	6492.69
6	AM	255.4
1	AM	2789.03
2	AP	5580.18
5	AP	2929.67

only showing top 20 rows

Insights on Orders

```

1 total_order_df = ecommerce_df.agg(count("id").alias("total_orders"))
2 total_order_df.show()
3
4 ecommerce_datasets["total_orders"] = total_order_df

```

✓ - Command executed in 512 ms on 8:50:52 AM, 9/05/24

total_orders
100000

Orders by Day

```

1 total_order_by_day_df = ecommerce_df.groupBy("day_of_week").agg(count("id").alias("total_orders")).orderBy("total_orders")
2 total_order_by_day_df.show()
3
4 ecommerce_datasets["total orders by day"] = total_order_by_day_df

```

day_of_week	total_orders
Sunday	10944
Monday	12034
Saturday	14199
Friday	14857
Thursday	15634
Wednesday	16045
Tuesday	16287

```

1 total_order_by_day_and_city_df = ecommerce_df.groupBy("day_of_week", "customer_city").agg(count("id").alias("total_orders")).orderBy("customer_city")
2 total_order_by_day_and_city_df.show()
3
4 ecommerce_datasets["total_orders_by_day_and_city"] = total_order_by_day_and_city_df

```

59 | ✓ - Command executed in 1 sec 163 ms on 8:50:55 AM, 9/05/24

```

+-----+-----+-----+
|day_of_week|customer_city|total_orders|
+-----+-----+-----+
| Friday|ALMIRANTE TAMANDA...|1|
| Tuesday|ALMIRANTE TAMANDA...|1|
| Sunday|ALTA FLORESTA D'O...|2|
| Tuesday|ALTO ALEGRE DOS P...|1|
| Sunday|ALTO ALEGRE DOS P...|1|
| Monday|ALTO ALEGRE DOS P...|1|
| Friday|ALVORADA D'OESTE|1|
| Tuesday|ALVORADA D'OESTE|1|
| Saturday|Abadia dos Dourados|1|
| Tuesday|Abadia dos Dourados|2|
| Tuesday|Abadiania|1|
| Friday|Abadiania|1|
| Thursday|Abaete|3|
| Tuesday|Abaete|1|
| Wednesday|Abaete|3|
| Friday|Abaete|1|
| Wednesday|Abaetetuba|3|
| Monday|Abaetetuba|2|
| Saturday|Abaetetuba|3|
| Friday|Abaetetuba|2|
+-----+-----+-----+
only showing top 20 rows

```

```

1 total_order_by_day_and_state_df = ecommerce_df.groupBy("day_of_week","customer_state").agg(count("id").alias("total_orders")).orderBy("customer_state")
2 total_order_by_day_and_state_df.show()
3
4 ecommerce_datasets["total_orders_by_day_and_state"] = total_order_by_day_and_state_df

```

```

+-----+-----+-----+
|day_of_week|customer_state|total_orders|
+-----+-----+-----+
| Thursday|AC|10|
| Monday|AC|7|
| Sunday|AC|11|
| Friday|AC|15|
| Saturday|AC|16|
| Wednesday|AC|13|
| Tuesday|AC|12|
| Saturday|AL|53|
| Sunday|AL|54|
| Monday|AL|59|
| Tuesday|AL|80|
| Wednesday|AL|52|
| Friday|AL|71|
| Thursday|AL|65|
| Sunday|AM|21|
| Friday|AM|12|
| Wednesday|AM|22|
| Saturday|AM|21|
| Monday|AM|20|
| Thursday|AM|28|
+-----+-----+-----+
only showing top 20 rows

```

Orders by Week

```

1 total_order_by_week_df = ecommerce_df.groupBy("week_of_month").agg(count("id").alias("total_orders")).orderBy("week_of_month")
2 total_order_by_week_df.show()
3
4 ecommerce_datasets["total_orders_by_week"] = total_order_by_week_df

```

✓ - Command executed in 1 sec 104 ms on 8:50:57 AM, 9/05/24

```

+-----+-----+
|week_of_month|total_orders|
+-----+-----+
|1|12202|
|2|23676|
|3|23255|
|4|23392|
|5|16160|
|6|1315|
+-----+-----+

```

```

1 total_order_by_week_and_city_df = ecommerce_df.groupBy("week_of_month","customer_city").agg(count("id").alias("total_orders")).orderBy("customer_city")
2 total_order_by_week_and_city_df.show()
3
4 ecommerce_datasets["total_orders_by_week_and_city"] = total_order_by_week_and_city_df

```

✓ - Command executed in 1 sec 85 ms on 8:50:59 AM, 9/05/24

```

|week_of_month|customer_city|total_orders|
+-----+-----+
|2|ALMIRANTE TAMANDA...|1|
|5|ALMIRANTE TAMANDA...|1|
|4|ALTA FLORESTA D'O...|1|
|1|ALTA FLORESTA D'O...|1|
|3|ALTO ALEGRE DOS P...|1|
|4|ALTO ALEGRE DOS P...|2|
|1|ALVORADA D'OESTE|1|
|5|ALVORADA D'OESTE|1|
|3|Abadia dos Dourados|3|
|5|Abadiania|1|
|1|Abadiania|1|
|5|Abaete|3|
|4|Abaete|2|
|2|Abaete|1|
|1|Abaete|1|
|3|Abaete|1|
|3|Abaetetuba|7|
|1|Abaetetuba|1|
|2|Abaetetuba|2|
|4|Abaetetuba|2|
+-----+-----+

```

only showing top 20 rows

```

1 total_order_by_week_and_state_df = ecommerce_df.groupBy("week_of_month","customer_state").agg(count("id").alias("total_orders")).orderBy("customer_state")
2 total_order_by_week_and_state_df.show()
3
4 ecommerce_datasets["total_orders_by_week_and_state"] = total_order_by_week_and_state_df

```

```

+-----+-----+-----+
|week_of_month|customer_state|total_orders|
+-----+-----+-----+
|          1|          AC|          15|
|          4|          AC|          20|
|          5|          AC|          15|
|          3|          AC|          20|
|          6|          AC|           1|
|          2|          AC|          13|
|          6|          AL|           4|
|          4|          AL|         104|
|          1|          AL|          52|
|          3|          AL|         113|
|          2|          AL|          89|
|          5|          AL|          72|
|          4|          AM|          41|
|          3|          AM|          30|
|          5|          AM|          18|
|          2|          AM|          41|
|          6|          AM|           2|
|          1|          AM|          22|
|          2|          AP|          25|
|          5|          AP|          14|
+-----+-----+-----+

```

only showing top 20 rows

Average of Misc. Columns

By Day

ML   

```

1 average_review_score_by_day_df = ecommerce_df.groupBy("day_of_week").agg(round(avg("review_score"),2).alias("average_review_score")).orderBy("average_review_score")
2 average_review_score_by_day_df.show()
3
4 ecommerce_datasets["average_review_score_by_day"] = average_review_score_by_day_df

```

✓ - Command executed in 1 sec 69 ms on 8:51:01 AM, 9/05/24

```

+-----+-----+
|day_of_week|average_review_score|
+-----+-----+
| Tuesday|          4.04|
| Friday|          4.04|
| Wednesday|          4.05|
| Saturday|          4.05|
| Sunday|          4.05|
| Thursday|          4.06|
| Monday|          4.06|
+-----+-----+

```

```

1 average_freight_value_by_day_df = ecommerce_df.groupBy("day_of_week").agg(round(avg("order_freight_value"),2).alias("average_freight_value"))\
2 .orderBy("average_freight_value")
3 average_freight_value_by_day_df.show()
4
5 ecommerce_datasets["average_freight_value_by_day"] = average_freight_value_by_day_df

```

```

+-----+-----+
|day_of_week|average_freight_value|
+-----+-----+
| Monday|          21.49|
| Tuesday|          21.5|
| Wednesday|          21.7|
| Saturday|          21.79|
| Thursday|          21.85|
| Friday|          21.92|
| Sunday|          21.98|
+-----+-----+

```

```

1 ecommerce_df = ecommerce_df.withColumn("time_to_approve_order",col("order_approved_at").cast("long") - col('order_purchase_timestamp').cast("long"))
2 ecommerce_df = ecommerce_df.withColumn("time_to_deliver_order",col("order_delivered_customer_date").cast("long") - col('order_purchase_timestamp').cast("long"))
3 display(ecommerce_df)

```


time_to_approve_order	time_to_deliver_order
660	728940
110580	1190760
1020	811680
1020	1141200
3720	248340
780	1429200
176580	176580
720	863100
116460	848340
600	1574340
540	1093020
720	497160
148020	1043940

```

1 average_time_to_approve_order_by_day_df = ecommerce_df.groupBy("day_of_week").agg(round(avg("time_to_approve_order"),2)\
2 .alias("average_time_to_approve_order")).orderBy("average_time_to_approve_order")
3 average_time_to_approve_order_by_day_df.show()
4
5 ecommerce_datasets["average_time_to_approve_order_by_day"] = average_time_to_approve_order_by_day_df

```

[67] ✓ - Command executed in 1 sec 183 ms on 8:51:05 AM, 9/05/24

```

+-----+-----+
|day_of_week|average_time_to_approve_order|
+-----+-----+
| Thursday|          31410.68|
| Friday|          31577.61|
| Wednesday|         32383.93|
| Tuesday|          33473.55|
| Monday|          38871.57|
| Saturday|         48351.09|
| Sunday|          52226.07|
+-----+-----+

```

```

1 average_time_to_deliver_order_by_day_df = ecommerce_df.groupBy("day_of_week").agg(round(avg("time_to_deliver_order"),2)\
2 .alias("average_time_to_approve_order")).orderBy("average_time_to_approve_order")
3 average_time_to_deliver_order_by_day_df.show()
4
5 ecommerce_datasets["average_time_to_deliver_order_by_day"] = average_time_to_deliver_order_by_day_df

```

```

+-----+-----+
|day_of_week|average_time_to_approve_order|
+-----+-----+
| Monday|         1033518.33|
| Tuesday|         1036488.73|
| Wednesday|        1043121.45|
| Thursday|        1075211.61|
| Friday|         1101193.49|
| Sunday|         1154172.92|
| Saturday|        1172419.75|
+-----+-----+

```

By Week

```

1 average_review_score_by_week_df = ecommerce_df.groupBy("week_of_month").agg(round(avg("review_score"),2).alias("average_review_score")).orderBy("average_review_score")
2 average_review_score_by_week_df.show()
3
4 ecommerce_datasets["average_review_score_by_week"] = average_review_score_by_week_df

```

```

... +-----+-----+
|week_of_month|average_review_score|
+-----+-----+
|          3|          4.04|
|          5|          4.04|
|          4|          4.05|
|          2|          4.06|
|          1|          4.07|
|          6|          4.07|
+-----+-----+

```

```

1 average_freight_value_by_week_df = ecommerce_df.groupBy("week_of_month").agg(round(avg("order_freight_value"),2)\
2 .alias("average_freight_value")).orderBy("average_freight_value")
3 average_freight_value_by_week_df.show()
4
5 ecommerce_datasets["average_freight_value_by_week"] = average_freight_value_by_week_df

```

[70] ✓ - Command executed in 1 sec 59 ms on 8:51:09 AM, 9/05/24

```

... +-----+-----+
|week_of_month|average_freight_value|
+-----+-----+
|          6|          21.43|
|          3|          21.62|
|          5|          21.71|
|          2|          21.78|
|          1|          21.82|
|          4|          21.82|
+-----+-----+

```

```

1 average_time_to_approve_order_by_week_df = ecommerce_df.groupBy("week_of_month").agg(round(avg("time_to_approve_order"),2)\
2 .alias("average_time_to_approve_order")).orderBy("average_time_to_approve_order")
3 average_time_to_approve_order_by_week_df.show()
4
5 ecommerce_datasets["average_time_to_approve_order_by_week"] = average_time_to_approve_order_by_week_df

```

[71] ✓ - Command executed in 1 sec 34 ms on 8:51:10 AM, 9/05/24

```

... +-----+-----+
|week_of_month|average_time_to_approve_order|
+-----+-----+
|          3|          34514.17|
|          2|          37290.22|
|          5|          37842.18|
|          6|          37842.89|
|          4|          39152.78|
|          1|          40009.4|
+-----+-----+

```

```

1 average_time_to_deliver_order_by_week_df = ecommerce_df.groupBy("week_of_month").agg(round(avg("time_to_deliver_order"),2)\
2 .alias("average_time_to_deliver_order")).orderBy("average_time_to_deliver_order")
3 average_time_to_deliver_order_by_week_df.show()
4
5 ecommerce_datasets["average_time_to_deliver_order_by_week"] = average_time_to_deliver_order_by_week_df

```

```

... +-----+
| week_of_month | average_time_to_deliver_order |
+-----+
| 6 | 812527.32 |
| 5 | 1073336.96 |
| 3 | 1082696.05 |
| 2 | 1089760.94 |
| 4 | 1095327.44 |
| 1 | 1105146.08 |
+-----+

```

Extra Look at Freight Costs

```

1 total_freight_charges_df = ecommerce_df.agg(round(sum("order_freight_value"),2).alias("total_freight_charges"))
2 total_freight_charges_df.show()
3
4 ecommerce_datasets["total_freight_charges"] = total_freight_charges_df

```

[73] ✓ - Command executed in 507 ms on 8:51:12 AM, 9/05/24

```

+-----+
| total_freight_charges |
+-----+
| 2174127.92 |
+-----+

```

```

1 average_freight_charges_by_city_df = ecommerce_df.groupBy("customer_city").agg(round(avg("order_freight_value"),2).alias("average_freight_charges"))
2 average_freight_charges_by_city_df.show()
3
4 ecommerce_datasets["average_freight_charges_by_city"] = average_freight_charges_by_city_df

```

[74] ✓ - Command executed in 524 ms on 8:51:12 AM, 9/05/24

```

... +-----+
| customer_city | average_freight_charges |
+-----+
| Anaruama | 24.3 |
| Guidoal | 17.61 |
| Piranguinho | 16.18 |
| Tres Pontas | 22.49 |
| Senador Guilmard | 68.55 |
| Rio Novo | 15.56 |
| Carrancas | 16.09 |
| Fronteira | 20.64 |
| Utinga | 26.07 |
| Assis Brasil | 24.84 |
| Pianco | 65.21 |
| Macaubas | 45.09 |
| Livramento | 25.47 |
| Cristalândia | 48.81 |
| Machados | 42.51 |
| Rio do Campo | 25.14 |
| Purilandia | 27.49 |
| Alambari | 18.78 |
| Guajará-Mirim | 38.67 |
| Tapes | 21.69 |
+-----+

```

```
1 average_freight_charges_by_state_df = ecommerce_df.groupby("customer_state").agg(round(avg("order_freight_value"),2).alias("average_freight_charges"))
2 average_freight_charges_by_state_df.show()
3
4 ecommerce_datasets["average_freight_charges_by_state"] = average_freight_charges_by_state_df
```

[75] ✓ - Command executed in 509 ms on 8:51:13 AM, 9/05/24

```
+-----+-----+
|customer_state|average_freight_charges|
+-----+-----+
|SC|23.25|
|RO|51.16|
|PI|41.18|
|AM|36.32|
|RR|46.53|
|GO|24.44|
|TO|38.92|
|MT|29.8|
|SP|16.63|
|ES|23.97|
|PB|41.15|
|RS|23.64|
|MS|24.97|
|AL|37.71|
|MG|22.82|
|PA|37.95|
|BA|29.07|
|SE|39.24|
|PE|34.65|
|CE|35.23|
+-----+-----+
only showing top 20 rows
```

Week 4: Data analysis and visualization

Step 1: Write the results into HDFS

Exporting Data

Exporting to Data Lake

```
1 for i in ecommerce_datasets:
2     ecommerce_datasets[i].write.format("csv").option("header", "True").save('abfss://ecommercecontainerxg@comcercedatag.dfs.core.windows.net/results/{}'.format(i,'n'))
```

[57] ✓ - Command executed in 1 min 17 sec 562 ms on 1:42:22 PM, 9/04/24

← → ↕ ecommercecontainerxg > results

Name	Last Modified	Content Type	Size
average_freight_charges_by_city	9/4/2024, 1:42:15 PM	Folder	
average_freight_charges_by_state	9/4/2024, 1:42:17 PM	Folder	
average_freight_value_by_day	9/4/2024, 1:41:52 PM	Folder	
average_freight_value_by_week	9/4/2024, 1:42:03 PM	Folder	
average_review_score_by_day	9/4/2024, 1:41:49 PM	Folder	
average_review_score_by_week	9/4/2024, 1:42:00 PM	Folder	
average_time_to_approve_order_by_day	9/4/2024, 1:41:54 PM	Folder	
average_time_to_approve_order_by_week	9/4/2024, 1:42:06 PM	Folder	
average_time_to_deliver_order_by_day	9/4/2024, 1:41:57 PM	Folder	
average_time_to_deliver_order_by_week	9/4/2024, 1:42:09 PM	Folder	
total_freight_charges	9/4/2024, 1:42:12 PM	Folder	
total_orders	9/4/2024, 1:41:28 PM	Folder	
total_orders_by_day	9/4/2024, 1:41:31 PM	Folder	
total_orders_by_day_and_city	9/4/2024, 1:41:34 PM	Folder	
total_orders_by_day_and_state	9/4/2024, 1:41:37 PM	Folder	
total_orders_by_week	9/4/2024, 1:41:40 PM	Folder	
total_orders_by_week_and_city	9/4/2024, 1:41:43 PM	Folder	
total_orders_by_week_and_state	9/4/2024, 1:41:46 PM	Folder	
total_sales	9/4/2024, 1:41:05 PM	Folder	

Showing 1 to 25 of 25 cached items

Step 2: Save the final dataset into object storage service per the cloud platform

Integrate

Filter resources by name

Pipelines

BlobtoLakePL

LakeResultstoBlob

Activities

Search activities

Synapse

Move and transform

Azure Data Explorer

Azure Function

Batch Service

Databricks

ecommerce-blob-data

LakeResultstoBlob

ecommerce_data

ecommercecontainer...

Validate

Validate copy runtime

Debug

Add trigger

Copy data

CopyLakeResultstoBlob

General

Source

Sink

Mapping

Settings

User properties

Source dataset *

ecommerceResultsCSVLakeDS

Open

New

Preview data

Learn more

File path type

File path in dataset

Wildcard file path

List of files

Wildcard paths

ecommercecontainerxg / results / *

Filter by last modified

Start time (UTC)

End time (UTC)

Recursively

Enable partitions discovery

Max concurrent connections

General

Source

Sink

Mapping

Settings

User properties

Sink dataset *

ecommerceResultsCSVBlobDS

Open

New

Learn more

Copy behavior

Preserve hierarchy

Max concurrent connections

Block size (MB)

Metadata

New

Quote all text

Pipeline run ID: 96e159ed-2709-4174-b55f-9e3986c65cb9

Pipeline status: Succeeded

View debug run consumption

All status

Monitor in Azure Metrics

Export to CSV

Showing 1 - 1 of 1 items

Activity name	Activity status	Activity type	Run start	Duration	Integration runtime	Us
CopyLakeResultstoBlob	Succeeded	Copy data	9/5/2024, 12:40:20 PM	15s	AutoResolveIntegration	

← → ↕ ↑

ecommerce-blob-data

Search by prefix (case-sensitive)

Name	Access Tier	Access Tier Last Modified	Last Modified	Blob Type	Content Type	Size	Status	Remaining Days	Deleted Time	Lease State	Disk Name	VM Name
average_freight_charges_by_city					Folder							
average_freight_charges_by_state					Folder							
average_freight_value_by_day					Folder							
average_freight_value_by_week					Folder							
average_review_score_by_day					Folder							
average_review_score_by_week					Folder							
average_time_to_approve_order_by_day					Folder							
average_time_to_approve_order_by_week					Folder							
average_time_to_deliver_order_by_day					Folder							
average_time_to_deliver_order_by_week					Folder							
total_freight_charges					Folder							
total_orders					Folder							
total_orders_by_day					Folder							
total_orders_by_day_and_city					Folder							
total_orders_by_day_and_state					Folder							
total_orders_by_week					Folder							
total_orders_by_week_and_city					Folder							
total_orders_by_week_and_state					Folder							
total_sales					Folder							

Showing 1 to 26 of 26 cached items

Step 3: Create a DB cluster that is also a NoSQL using the relevant service on the cloud platform

Step 4: Save insights in the NoSQL DB mentioned in the previous step

Exporting to CosmosDB

```

1  for i in ecommerce_datasets:
2      #Making id column
3      dataset_x = ecommerce_datasets[i]
4      dataset_x = dataset_x.withColumn("index", monotonically_increasing_id())
5      dataset_x = dataset_x.withColumn("table_name", lit(i))
6      dataset_x = dataset_x.withColumn("id", concat(col("table_name"),lit("-"), col("index")))
7      dataset_x = dataset_x.drop("index")
8      #Exporting to CosmosDB
9      config = {
10         'spark.cosmos.accountEndpoint': 'https://cosmosdb-xg-sql.documents.azure.com:443/',
11         'spark.cosmos.accountKey': 'wVkkMa27deTM6hHemUN68KJHSeJmvI1ZJpYSQXaSZ6qFNURAAaIlyvOjutbwYizQFfaSXKSxwLvmIACDbcbmCdA==',
12         'spark.cosmos.database': 'synapselinkdb',
13         'spark.cosmos.container': 'results',
14     }
15     dataset_x.write.format("cosmos.oltp").options(**config).mode('append').save()

```

[81] ✓ - Command executed in 6 min 26 sec 826 ms on 9:03:33 AM, 9/05/24

...

Azure Cosmos DB account

Search

- Overview
- Activity log
- Access control (IAM)
- Tags
- Diagnose and solve problems
- Cost Management
- Quick start
- Data Explorer
 - Settings
 - Integrations
 - Containers
 - Monitoring
 - Automation
 - Help

Query with AI using Microsoft Copilot for Azure in Cosmos DB! Ask your subscription admin to enable the preview today. [Learn more](#) [Enroll](#)

Execute Query Save Query Download Query View

+ New Container

Home

synapselinkdb

Scale

product_arr

results

Home results.items resultsQuery 1 x

1 SELECT distinct(c.table_name) FROM c

Results Query Stats

1 - 25

```
{
  "table_name": "average_freight_charges_by_city"
},
{
  "table_name": "average_freight_charges_by_state"
},
{
  "table_name": "average_freight_value_by_day"
},
{
  "table_name": "average_freight_value_by_week"
},
{
  "table_name": "average_review_score_by_day"
},
{
  "table_name": "average_review_score_by_week"
},
{
  "table_name": "average_time_to_approve_order_by_day"
},
{
  "table_name": "average_time_to_approve_order_by_week"
},
{
  "table_name": "average_time_to_deliver_order_by_day"
},
{
  "table_name": "average_time_to_deliver_order_by_week"
}
```

+ New Container

Home

SELECT * FROM c

Home	<input checked="" type="checkbox"/>	id	/table_name
synapselinkdb	<input checked="" type="checkbox"/>	total_sales-0	total_sales
Scale	<input type="checkbox"/>	total_sales_by_day-0	total_sales_by_day
product_arr	<input type="checkbox"/>	total_sales_by_day-1	total_sales_by_day
results	<input type="checkbox"/>	total_sales_by_day-2	total_sales_by_day
Items	<input type="checkbox"/>	total_sales_by_day-3	total_sales_by_day
Settings	<input type="checkbox"/>	total_sales_by_day-4	total_sales_by_day
Stored Procedures	<input type="checkbox"/>	total_sales_by_day-5	total_sales_by_day
User Defined Functions	<input type="checkbox"/>	total_sales_by_day-6	total_sales_by_day
Triggers	<input type="checkbox"/>	total_sales_by_day_and_city-0	total_sales_by_day_and_city
	<input type="checkbox"/>	total_sales_by_day_and_city-1	total_sales_by_day_and_city
	<input type="checkbox"/>	total_sales_by_day_and_city-2	total_sales_by_day_and_city

```
1 {
2   "total_sales": 12841476.98,
3   "table_name": "total_sales",
4   "id": "total_sales-0",
5   "_rid": "UV15A94C4Y48AAAAAAAAAA==",
6   "_self": "dbs/UV15AA=/colls/UV15A94C4Y4=/docs/UV15A94C4Y48AAAAAAAAAA==/",
7   "_etag": "\"04002022-0000-0300-0000-66d9d4d00000\"",
8   "_attachments": "attachments/",
9   "_ts": 1725551832
10 }
```